

Performance Analysis of ISAs

The 4 Instruction Set Architectures that we were asked to compare are very similar in terms of instructions related to code flow (branching, goto, etc...). However, because of their different makeup when it comes to arithmetic operations, there are some serious differences in how the different ISAs perform both in memory access, ability to perform optimizations, and code size. The compilers that I wrote in a lot of cases demonstrate these innate properties related to their target ISA but due to lack of optimizations some ISAs (particularly the load and store architecture) do not match the expected results.

CPI performance of each ISA

Lets first look at the raw performance of each ISA. For this performance metric I used the average CPI x the instruction count to get the total cost of each program.

The CPI list for the following metrics was :

ADD,10
SUB,10
MULT,10
DIV,10
PUSH,20
POP,20
STORE,40
LOAD,35
GOTO,5
BLT,8
BGT,8
BLE,9
BGE,9
BEQ,5
BNE,5
CLEAR,2

Comparing a memory intensive (medium) program, I observed the following data:

ISA	Average CPI	Total Cost
Memory Memory	37	3265
Stack	19	3575
Load Store	36	6565
Accumulator	17	4485

This shows that we the least cpu cost comes from the Memory Memory ISA and the Stack ISA. This is consistent with what we would expect as they would have the least amount of instructions. Interestingly, while the Load and Store ISA had a similar low instruction count to the Stack ISA, it's

average CPI eventually made it have the highest total cost.

When we compare a cpu intensive (medium) program, we get slightly different results:

ISA	Average CPI	Total Cost
Memory Memory	11	815
Stack	15	4999
Load Store	34	7215
Accumulator	15	4387

In this simulation, we see that the Memory Memory ISA is the clear winner which is expected as it would have a lot less instructions to perform as there are no load or store instructions to perform. We also see the simplicity of the Accumulator and Stack ISAs show as well with their relatively low total cost. Again we see the Load and Store ISA come in last. This is a bit misleading however because the compiler is not doing any code optimizations. Theoretically, the Load and Store ISAs cost should be much lower as many load and store instructions would be removed from the optimized assembly. This could be accomplished using dynamic register allocation, but as that wasn't necessarily part of the assignment, it didn't get included in the compiler. In fact the Load and Store compiler is completely un-optimized to the point that it only uses at most 3 registers and always reloads data fresh into each needed register.

Memory Usage of each ISA

Another metric we can look at is the memory usage and memory access of each ISA. For these numbers we're again using the same CPI and the memory intensive (medium) program.

ISA	Memory Access Count	Memory Access (in bytes)
Memory Memory	178	712
Stack	178	712
Load Store	178	712
Accumulator	178	712

Interestingly, despite the fact that the numbers look wrong, this is almost what we should expect to see. Memory Memory, Stack, and Accumulator ISAs all should have the same memory access usage. The one erroneous number is the Load and Store ISA memory usage. This is because again, the Load and Store compiler target is horribly un-optimized. With dynamic register allocation we could be looking at a much smaller usage.

With the cpu intensive program, we see the same results as the memory intensive program (table omitted) .

Code size of each ISA

Another important aspect to look at is the generated code size. The calculated code sizes were consistent with what we'd expect to see theoretically.

For the memory intensive (medium) program, the following data was observed.

ISA	Instructions	Code Size (in bytes)
Memory Memory	88	444
Stack	188	544
Load Store	178	884
Accumulator	263	619

We expect that the Memory Memory ISA target will generate the smallest number of instructions and code because it doesn't have to load/store memory in cases of arithmetic instructions. The Load and Store ISA does remarkably well in the instruction count department considering that it is horribly un-optimized. However because each instruction almost always has 1-3 memory addresses associated with it, it's total code size balloons up to be the worst option.

With the cpu intensive program, we see much the same results.

ISA	Instructions	Code Size (in bytes)
Memory Memory	74	500
Stack	333	759
Load Store	210	1076
Accumulator	284	710

The difference here is that stack based arithmetic operations require more instructions to perform because of the extra push/pop commands that are required to do arithmetic operations.

Memory Bandwidth of each ISA

Lastly we can look at the total memory bandwidth that is used for each ISA. The total memory bandwidth is the combination of the code size and the memory access bandwidth.

For the memory intensive program, the following was observed.

ISA	Total Memory Bandwidth
Memory Memory	1156
Stack	1256
Load Store	1596

Accumulator	1331
-------------	------

As can be seen, the Load and Store ISA compiler generates the most memory usage of the different ISAs. This is very erroneous as this compiler target is very unoptimized. Theoretically we'd see the Load and Store ISA having the best memory bandwidth (which is why it tends to be used in today's CPU hardware).

For the CPU intensive program we get much the same erroneous picture.

ISA	Total Memory Bandwidth
Memory Memory	1276
Stack	1535
Load Store	1852
Accumulator	1486

Again, this result is skewed by the fact that the Load and Store compiler doesn't have any assembly level optimizations.

Summary

Overall each ISA provides advantages and disadvantages in theoretical and real world use. The Memory Memory ISA produces dense code with low CPU cost. However it is hard to implement in hardware in a cost effective way. Stack and Accumulator ISAs provide a simple to use ISA with good memory usage and low CPI, but they can prove to be costly when it comes to instruction counts. Also they tend to be hard to optimize for because resources must be stored in the CPU and cannot be pipelined easily. Lastly the Load and Store ISA certainly seems to be a bad choice when it comes to ISAs, but both theoretically and in practice, the Load and Store ISA allows great optimization gains (when properly optimized and via pipelining) and allows simpler (ie cheaper) CPUs to be manufactured.